

ML0101EN-Clus-DBSCN-weather-py-v1

March 4, 2019

Density-Based Clustering

Most of the traditional clustering techniques, such as k-means, hierarchical and fuzzy clustering, can be used to group data without supervision.

However, when applied to tasks with arbitrary shape clusters, or clusters within cluster, the traditional techniques might be unable to achieve good results. That is, elements in the same cluster might not share enough similarity or the performance may be poor. Additionally, Density-based Clustering locates regions of high density that are separated from one another by regions of low density. Density, in this context, is defined as the number of points within a specified radius.

In this section, the main focus will be manipulating the data and properties of DBSCAN and observing the resulting clustering.

Table of contents

Clustering with Randomly Generated Data

Data generation

Modeling

Distinguishing Outliers

Data Visualization

Weather Station Clustering with DBSCAN & scikit-learn

Loading data

Overview data

Data cleaning

Data selection

Clustering

Visualization of clusters based on location

Clustering of stations based on temperature

Visualization of clusters based on temperature

Import the following libraries:

 numpy as np

 DBSCAN from sklearn.cluster

```

<li> <b>make_blobs</b> from <b>sklearn.datasets.samples_generator</b> </li>
<li> <b>StandardScaler</b> from <b>sklearn.preprocessing</b> </li>
<li> <b>matplotlib.pyplot</b> as plt </li>

```

Remember %matplotlib inline to display plots

```

In [ ]: # Notice: For visualization of map, you need basemap package.
        # if you dont have basemap install on your machine, you can use the following line to install
        # !conda install -c conda-forge basemap=1.1.0 matplotlib=2.2.2 -y
        # Notice: you might have to refresh your page and re-run the notebook after installation

```

```

In [1]: import numpy as np
        from sklearn.cluster import DBSCAN
        from sklearn.datasets.samples_generator import make_blobs
        from sklearn.preprocessing import StandardScaler
        import matplotlib.pyplot as plt
        %matplotlib inline

```

Data generation

The function below will generate the data points and requires these inputs:

```

<li> <b>centroidLocation</b>: Coordinates of the centroids that will generate the random data. </li>
<ul> <li> Example: input: [[4,3], [2,-1], [-1,4]] </li> </ul>
<li> <b>numSamples</b>: The number of data points we want generated, split over the number of centroids. </li>
<ul> <li> Example: 1500 </li> </ul>
<li> <b>clusterDeviation</b>: The standard deviation between the clusters. The larger the number, the more spread out the data points will be. </li>
<ul> <li> Example: 0.5 </li> </ul>

```

```

In [2]: def createDataPoints(centroidLocation, numSamples, clusterDeviation):
        # Create random data and store in feature matrix X and response vector y.
        X, y = make_blobs(n_samples=numSamples, centers=centroidLocation,
                           cluster_std=clusterDeviation)

        # Standardize features by removing the mean and scaling to unit variance
        X = StandardScaler().fit_transform(X)
        return X, y

```

Use createDataPoints with the 3 inputs and store the output into variables X and y.

```

In [4]: X, y = createDataPoints([[4,3], [2,-1], [-1,4]] , 1500, 0.5)
        X,y

```

```

Out[4]: (array([[-1.28973488,  0.69863618],
                [ 0.80764974,  0.62753264],
                [ 1.19524154,  0.39901638],
                ...,
                [ 0.05757003, -1.04586961],
                [ 1.17998356,  0.46101371],
                [ 0.77687284,  0.77710061]]), array([2, 0, 0, ..., 1, 0, 0]))

```

Modeling

DBSCAN stands for Density-Based Spatial Clustering of Applications with Noise. This technique is one of the most common clustering algorithms which works based on density of object. The whole idea is that if a particular point belongs to a cluster, it should be near to lots of other points in that cluster.

It works based on two parameters: Epsilon and Minimum Points

Epsilon determine a specified radius that if includes enough number of points within, we call it dense area

minimumSamples determine the minimum number of data points we want in a neighborhood to define a cluster.

```
In [60]: epsilon = 0.3
         minimumSamples = 7

         db = DBSCAN(eps=epsilon, min_samples=minimumSamples).fit(X)
         labels = db.labels_
```

Distinguishing Outliers

Lets Replace all elements with 'True' in core_samples_mask that are in the cluster, 'False' if the points are outliers.

```
In [10]: # First, create an array of booleans using the labels from db.
         core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
         core_samples_mask[db.core_sample_indices_] = True
         core_samples_mask
```

```
Out[10]: array([ True,  True,  True, ...,  True,  True,  True])
```

```
In [16]: # Number of clusters in labels, ignoring noise if present.
         n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)
         n_clusters_
```

```
Out[16]: 3
```

```
In [32]: # Remove repetition in labels by turning it into a set.
         unique_labels = set(labels)
         unique_labels
```

```
Out[32]: {-1, 0, 1, 2}
```

Data visualization

```
In [71]: # Create colors for the clusters.
         colors = plt.cm.Spectral(np.linspace(0, 1, len(unique_labels)))
         colors
```

```
Out[71]: array([[0.61960784, 0.00392157, 0.25882353, 1.          ],
                [0.99346405, 0.74771242, 0.43529412, 1.          ],
                [0.74771242, 0.89803922, 0.62745098, 1.          ],
                [0.36862745, 0.30980392, 0.63529412, 1.          ]])
```

```

In [72]: # Plot the points with colors
for k, col in zip(unique_labels, colors):
    if k == -1:
        # Black used for noise.
        col = 'k'

    class_member_mask = (labels == k)

    # Plot the datapoints that are clustered
    xy = X[class_member_mask & core_samples_mask]
    plt.scatter(xy[:, 0], xy[:, 1], s=50, c=col, marker='o', alpha=0.5)

    # Plot the outliers
    xy = X[class_member_mask & ~core_samples_mask]
    plt.scatter(xy[:, 0], xy[:, 1], s=50, c=col, marker='o', alpha=0.5)

```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-


'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-




0.1 Practice

To better understand differences between partitional and density-based clustering, try to cluster the above dataset into 3 clusters using k-Means.

Notice: do not generate data again, use the same dataset as above.

```
In [37]: # write your code here
from sklearn.cluster import KMeans
k_means3 = KMeans(init = "k-means++", n_clusters = 3, n_init = 12)
k_means3.fit(X)
fig = plt.figure(figsize=(6, 4))
colors = plt.cm.Spectral(np.linspace(0, 1, len(set(k_means3.labels_))))
ax = fig.add_subplot(1, 1, 1)
for k, col in zip(range(len(k_means3.cluster_centers_)), colors):
    my_members = (k_means3.labels_ == k)
    cluster_center = k_means3.cluster_centers_[k]
    ax.plot(X[my_members, 0], X[my_members, 1], 'w', markerfacecolor=col, marker='.')
    ax.plot(cluster_center[0], cluster_center[1], 'o', markerfacecolor=col, markeredgecolor='k')
plt.show()
```



Double-click [here](#) for the solution.

Weather Station Clustering using DBSCAN & scikit-learn

DBSCAN is specially very good for tasks like class identification on a spatial context. The wonderful attribute of DBSCAN algorithm is that it can find out any arbitrary shape cluster without getting affected by noise. For example, this following example cluster the location of weather stations in Canada. DBSCAN can be used here, for instance, to find the group of stations which show

the same weather condition. As you can see, it not only finds different arbitrary shaped clusters, can find the denser part of data-centered samples by ignoring less-dense areas or noises.

let's start playing with the data. We will be working according to the following workflow:

0.1.1 About the dataset

Environment Canada

Monthly Values for July - 2015

<th>Name in the table</th>

<th>Meaning</th>

<td>Stn_Name</td>

<td>Station Name</td>

<td>Lat</td>

<td>Latitude (North+, degrees)</td>

<td>Long</td>

<td>Longitude (West - , degrees)</td>

<td>Prov</td>

<td>Province</td>

<td>Tm</td>

<td>Mean Temperature (řC)</td>

<td>DwTm</td>

<td>Days without Valid Mean Temperature</td>

<td>D</td>

<td>Mean Temperature difference from Normal (1981-2010) (řC)</td>

<td>Tx</td>

<td>Highest Monthly Maximum Temperature (řC)</td>

<td>DwTx</td>

<td>Days without Valid Maximum Temperature</td>

<td>Tn</td>

<td>Lowest Monthly Minimum Temperature (řC)</td>

<td>DwTn</td>

<td>Days without Valid Minimum Temperature</td>

<td>S</td>

<td>Snowfall (cm)</td>

<td>DwS</td>

<td>Days without Valid Snowfall</td>

```

<td>S%N</td>
<td>Percent of Normal (1981-2010) Snowfall</td>

<td><font color = "green"><strong>P</font></td>
<td><font color = "green"><strong>Total Precipitation (mm)</font></td>

<td>DwP</td>
<td>Days without Valid Precipitation</td>

<td>P%N</td>
<td>Percent of Normal (1981-2010) Precipitation</td>

<td>S_G</td>
<td>Snow on the ground at the end of the month (cm)</td>

<td>Pd</td>
<td>Number of days with Precipitation 1.0 mm or more</td>

<td>BS</td>
<td>Bright Sunshine (hours)</td>

<td>DwBS</td>
<td>Days without Valid Bright Sunshine</td>

<td>BS%</td>
<td>Percent of Normal (1981-2010) Bright Sunshine</td>

<td>HDD</td>
<td>Degree Days below 18 řC</td>

<td>CDD</td>
<td>Degree Days above 18 řC</td>

<td>Stn_No</td>
<td>Climate station identifier (first 3 digits indicate    drainage basin, last 4 characters are

<td>NA</td>
<td>Not Available</td>

```

0.1.2 1-Download data

To download the data, we will use `!wget` to download it from IBM Object Storage.

Did you know? When it comes to Machine Learning, you will likely be working with large da

```

In [73]: !wget -O weather-stations20140101-20141231.csv https://s3-api.us-geo.objectstorage.soft
--2019-03-04 20:27:42--  https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/Cogni
Resolving s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net).
Connecting to s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.n
HTTP request sent, awaiting response... 200 OK

```

Length: 129821 (127K) [text/csv]

Saving to: weather-stations20140101-20141231.csv

weather-stations201 100%[=====>] 126.78K --.-KB/s in 0.06s

2019-03-04 20:27:42 (1.94 MB/s) - weather-stations20140101-20141231.csv saved [129821/129821]

0.1.3 2- Load the dataset

We will import the .csv then we creates the columns for year, month and day.

```
In [78]: import csv
import pandas as pd
import numpy as np

filename='weather-stations20140101-20141231.csv'

#Read csv
pdf = pd.read_csv(filename)
pdf.head(5)
```

```
Out [78]:
```

	Stn_Name	Lat	Long	Prov	Tm	DwTm	D	Tx	DwTx	\
0	CHEMAINUS	48.935	-123.742	BC	8.2	0.0	NaN	13.5	0.0	
1	COWICHAN LAKE FORESTRY	48.824	-124.133	BC	7.0	0.0	3.0	15.0	0.0	
2	LAKE COWICHAN	48.829	-124.052	BC	6.8	13.0	2.8	16.0	9.0	
3	DISCOVERY ISLAND	48.425	-123.226	BC	NaN	NaN	NaN	12.5	0.0	
4	DUNCAN KELVIN CREEK	48.735	-123.728	BC	7.7	2.0	3.4	14.5	2.0	

	Tn	...	DwP	P%N	S_G	Pd	BS	DwBS	BS%	HDD	CDD	Stn_No
0	1.0	...	0.0	NaN	0.0	12.0	NaN	NaN	NaN	273.3	0.0	1011500
1	-3.0	...	0.0	104.0	0.0	12.0	NaN	NaN	NaN	307.0	0.0	1012040
2	-2.5	...	9.0	NaN	NaN	11.0	NaN	NaN	NaN	168.1	0.0	1012055
3	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1012475
4	-1.0	...	2.0	NaN	NaN	11.0	NaN	NaN	NaN	267.7	0.0	1012573

[5 rows x 25 columns]

0.1.4 3-Cleaning

Lets remove rows that don't have any value in the Tm field.

```
In [85]: pdf = pdf[pd.notnull(pdf["Tm"])]
pdf = pdf.reset_index(drop=True)
pdf.head(5)
```

```
Out [85]:
```

	Stn_Name	Lat	Long	Prov	Tm	DwTm	D	Tx	DwTx	\
0	CHEMAINUS	48.935	-123.742	BC	8.2	0.0	NaN	13.5	0.0	

1	COWICHAN LAKE FORESTRY	48.824	-124.133	BC	7.0	0.0	3.0	15.0	0.0
2	LAKE COWICHAN	48.829	-124.052	BC	6.8	13.0	2.8	16.0	9.0
3	DUNCAN KELVIN CREEK	48.735	-123.728	BC	7.7	2.0	3.4	14.5	2.0
4	ESQUIMALT HARBOUR	48.432	-123.439	BC	8.8	0.0	NaN	13.1	0.0

	Tn	...	S_G	Pd	BS	DwBS	BS%	HDD	CDD	Stn_No	xm	\
0	1.0	...	0.0	12.0	NaN	NaN	NaN	273.3	0.0	1011500	1.807806e+06	
1	-3.0	...	0.0	12.0	NaN	NaN	NaN	307.0	0.0	1012040	1.764329e+06	
2	-2.5	...	NaN	11.0	NaN	NaN	NaN	168.1	0.0	1012055	1.773336e+06	
3	-1.0	...	NaN	11.0	NaN	NaN	NaN	267.7	0.0	1012573	1.809363e+06	
4	1.9	...	NaN	12.0	NaN	NaN	NaN	258.6	0.0	1012710	1.841498e+06	

	ym
0	1.396332e+06
1	1.377564e+06
2	1.378409e+06
3	1.362546e+06
4	1.311615e+06

[5 rows x 27 columns]

0.1.5 4-Visualization

Visualization of stations on map using basemap package. The matplotlib basemap toolkit is a library for plotting 2D data on maps in Python. Basemap does not do any plotting on it's own, but provides the facilities to transform coordinates to a map projections.

Please notice that the size of each data points represents the average of maximum temperature for each station in a year.

```
In [83]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

llon=-140
ulon=-50
llat=40
ulat=65

pdf = pdf[(pdf['Long'] > llon) & (pdf['Long'] < ulon) & (pdf['Lat'] > llat) & (pdf['Lat'] < ulat)]

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and latitude (urcrnrlat)

my_map.drawcoastlines()
```

```

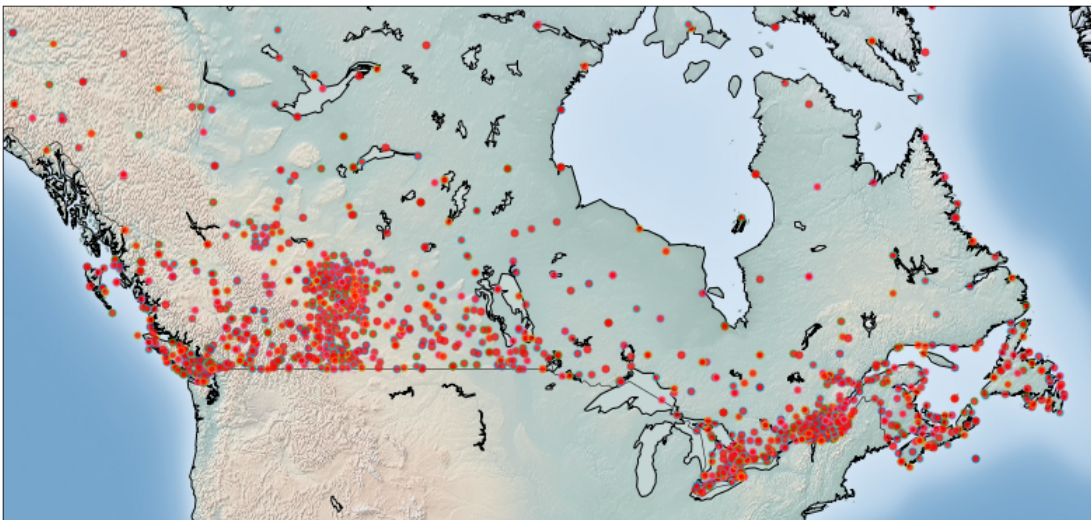
my_map.drawcountries()
# my_map.draumapboundary()
my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()

# To collect data based on stations

xs,ys = my_map(np.asarray(pdf.Long), np.asarray(pdf.Lat))
pdf['xm']= xs.tolist()
pdf['ym'] =ys.tolist()

#Visualization1
for index,row in pdf.iterrows():
#   x,y = my_map(row.Long, row.Lat)
    my_map.plot(row.xm, row.ym,markerfacecolor =([1,0,0]), marker='o', markersize= 5, a
#plt.text(x,y,stn)
plt.show()

```



Out[83]: <generator object DataFrame.iterrows at 0x7fdf9e3ba620>

0.1.6 5- Clustering of stations based on their location i.e. Lat & Lon

DBSCAN form sklearn library can runs DBSCAN clustering from vector array or distance matrix. In our case, we pass it the Numpy array Clus_dataSet to find core samples of high density and ex

```

In [86]: from sklearn.cluster import DBSCAN
import sklearn.utils
from sklearn.preprocessing import StandardScaler
sklearn.utils.check_random_state(1000)

```

```

Clus_dataSet = pdf[['xm', 'ym']]
Clus_dataSet = np.nan_to_num(Clus_dataSet)
Clus_dataSet = StandardScaler().fit_transform(Clus_dataSet)

# Compute DBSCAN
db = DBSCAN(eps=0.15, min_samples=10).fit(Clus_dataSet)
core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True
labels = db.labels_
pdf["Clus_Db"]=labels

realClusterNum=len(set(labels)) - (1 if -1 in labels else 0)
clusterNum = len(set(labels))

# A sample of clusters
pdf[["Stn_Name", "Tx", "Tm", "Clus_Db"]].head(5)

```

```

Out[86]:

```

	Stn_Name	Tx	Tm	Clus_Db
0	CHEMAINUS	13.5	8.2	0
1	COWICHAN LAKE FORESTRY	15.0	7.0	0
2	LAKE COWICHAN	16.0	6.8	0
3	DUNCAN KELVIN CREEK	14.5	7.7	0
4	ESQUIMALT HARBOUR	13.1	8.8	0

As you can see for outliers, the cluster label is -1

```

In [92]: set(labels)

```

```

Out[92]: {-1, 0, 1, 2, 3, 4}

```

0.1.7 6- Visualization of clusters based on location

Now, we can visualize the clusters using basemap:

```

In [88]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and latitude (urcrnrlat)

my_map.drawcoastlines()
my_map.drawcountries()
#my_map.drawmapboundary()

```

```

my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()

# To create a color map
colors = plt.get_cmap('jet')(np.linspace(0.0, 1.0, clusterNum))

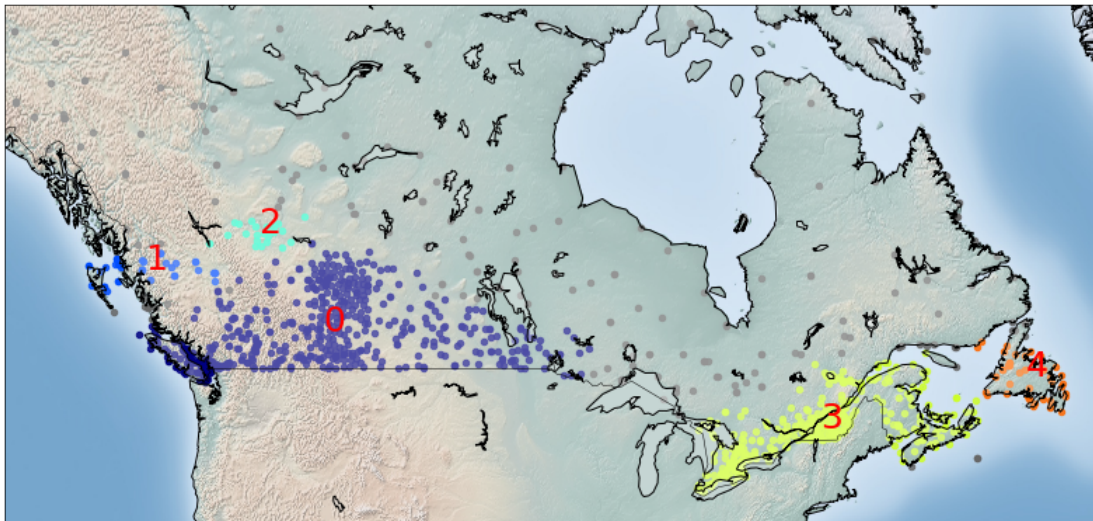
#Visualization1
for clust_number in set(labels):
    c=([0.4,0.4,0.4]) if clust_number == -1 else colors[np.int(clust_number)]
    clust_set = pdf[pdf.Clus_Db == clust_number]
    my_map.scatter(clust_set.xm, clust_set.ym, color=c, marker='o', s= 20, alpha = 0.5)
    if clust_number != -1:
        cenx=np.mean(clust_set.xm)
        ceny=np.mean(clust_set.ym)
        plt.text(cenx,ceny,str(clust_number), fontsize=25, color='red',)
        print ("Cluster "+str(clust_number)+', Avg Temp: '+ str(np.mean(clust_set.Tm)))

```

```

Cluster 0, Avg Temp: -5.538747553816046
Cluster 1, Avg Temp: 1.9526315789473685
Cluster 2, Avg Temp: -9.195652173913045
Cluster 3, Avg Temp: -15.300833333333333
Cluster 4, Avg Temp: -7.769047619047619

```



0.1.8 7- Clustering of stations based on their location, mean, max, and min Temperature

In this section we re-run DBSCAN, but this time on a 5-dimensional dataset:

```

In [93]: from sklearn.cluster import DBSCAN
import sklearn.utils
from sklearn.preprocessing import StandardScaler
sklearn.utils.check_random_state(1000)
Clus_dataSet = pdf[['xm', 'ym', 'Tx', 'Tm', 'Tn']]
Clus_dataSet = np.nan_to_num(Clus_dataSet)
Clus_dataSet = StandardScaler().fit_transform(Clus_dataSet)

# Compute DBSCAN
db = DBSCAN(eps=0.3, min_samples=10).fit(Clus_dataSet)
core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True
labels = db.labels_
pdf["Clus_Db"]=labels

realClusterNum=len(set(labels)) - (1 if -1 in labels else 0)
clusterNum = len(set(labels))

# A sample of clusters
pdf[["Stn_Name", "Tx", "Tm", "Clus_Db"]].head(5)

```

```

Out[93]:

```

	Stn_Name	Tx	Tm	Clus_Db
0	CHEMAINUS	13.5	8.2	0
1	COWICHAN LAKE FORESTRY	15.0	7.0	0
2	LAKE COWICHAN	16.0	6.8	0
3	DUNCAN KELVIN CREEK	14.5	7.7	0
4	ESQUIMALT HARBOUR	13.1	8.8	0

0.1.9 8- Visualization of clusters based on location and Temperature

```

In [94]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and latitude (urcrnrlat)

my_map.drawcoastlines()
my_map.drawcountries()
#my_map.drawmapboundary()
my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()

```

```

# To create a color map
colors = plt.get_cmap('jet')(np.linspace(0.0, 1.0, clusterNum))

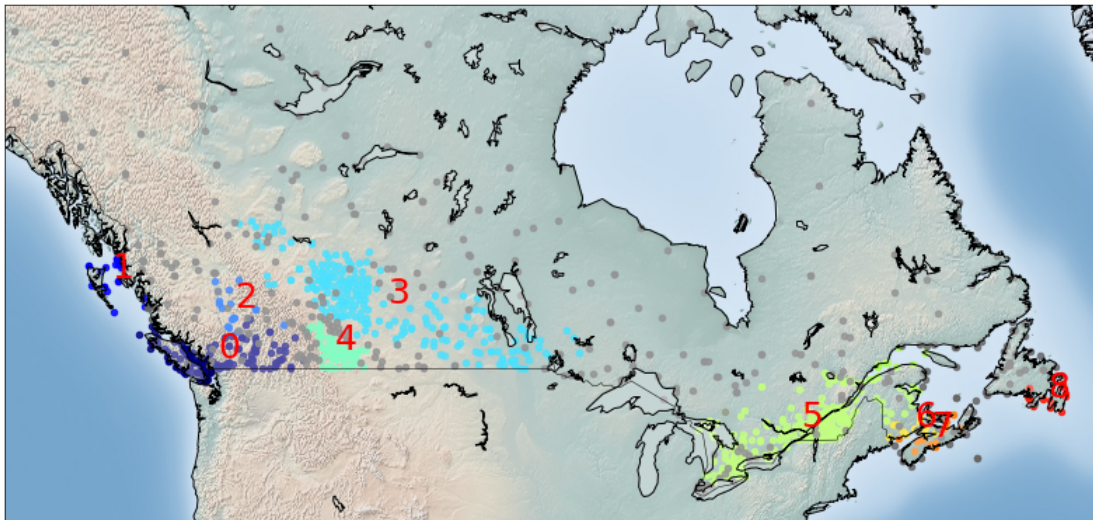
#Visualization1
for clust_number in set(labels):
    c=([0.4,0.4,0.4]) if clust_number == -1 else colors[np.int(clust_number)]
    clust_set = pdf[pdf.Clus_Db == clust_number]
    my_map.scatter(clust_set.xm, clust_set.ym, color=c, marker='o', s= 20, alpha = 0.
    if clust_number != -1:
        cenx=np.mean(clust_set.xm)
        ceny=np.mean(clust_set.ym)
        plt.text(cenx,ceny,str(clust_number), fontsize=25, color='red',)
        print ("Cluster "+str(clust_number)+', Avg Temp: ' + str(np.mean(clust_set.Tm)))

```

```

Cluster 0, Avg Temp: 6.221192052980132
Cluster 1, Avg Temp: 6.7900000000000001
Cluster 2, Avg Temp: -0.49411764705882344
Cluster 3, Avg Temp: -13.87720930232558
Cluster 4, Avg Temp: -4.186274509803922
Cluster 5, Avg Temp: -16.301503759398496
Cluster 6, Avg Temp: -13.599999999999998
Cluster 7, Avg Temp: -9.753333333333334
Cluster 8, Avg Temp: -4.258333333333334

```



Want to learn more?

IBM SPSS Modeler is a comprehensive analytics platform that has many machine learning algorithms. It has been designed to bring predictive intelligence to decisions made by individuals,

by groups, by systems – by your enterprise as a whole. A free trial is available through this course, available here: SPSS Modeler

Also, you can use Watson Studio to run these notebooks faster with bigger datasets. Watson Studio is IBM's leading cloud solution for data scientists, built by data scientists. With Jupyter notebooks, RStudio, Apache Spark and popular libraries pre-packaged in the cloud, Watson Studio enables data scientists to collaborate on their projects without having to install anything. Join the fast-growing community of Watson Studio users today with a free account at Watson Studio

Thanks for completing this lesson!

Author: Saeed Aghabozorgi

Saeed Aghabozorgi, PhD is a Data Scientist in IBM with a track record of developing enterprise level applications that substantially increases clients' ability to turn data into actionable knowledge. He is a researcher in data mining field and expert in developing advanced analytic methods like machine learning and statistical modelling on large datasets.

Copyright © 2018 Cognitive Class. This notebook and its source code are released under the terms of the MIT License.